



X-Ray Classifier

Alex Cross



/github.com/alcr0996



/in/alcr0996



alexfourdata@gmail.com

Background

This dataset was obtained from the Stanford ML group, which worked with a group of board certified radiologists to prepare and label the data.

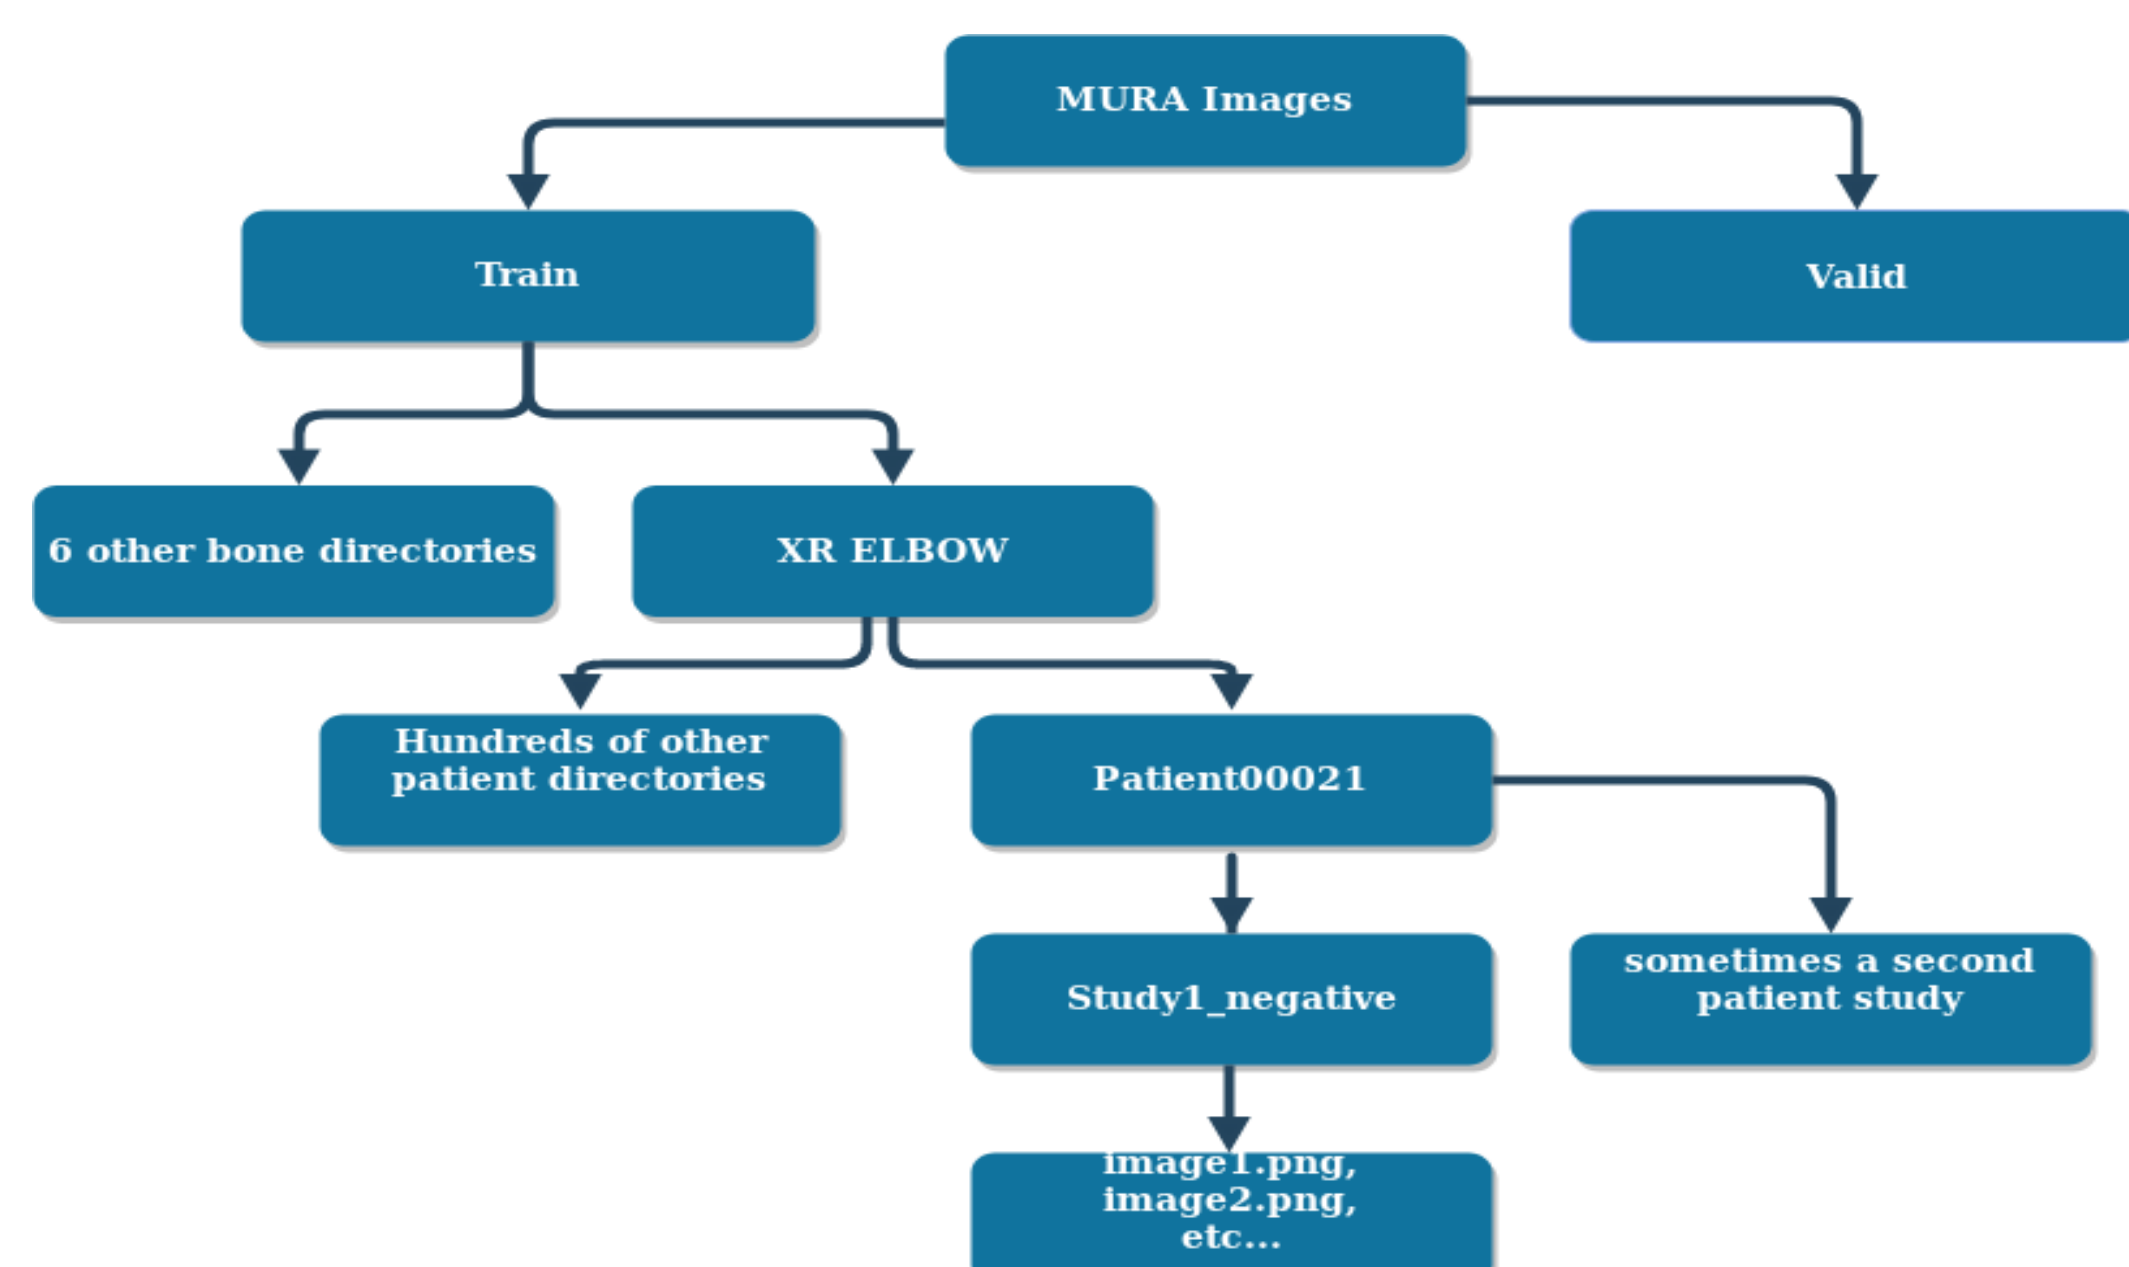
Creating a ML model that could assist in the identification of x-ray studies would be a major contribution to the healthcare industry.

Objectives

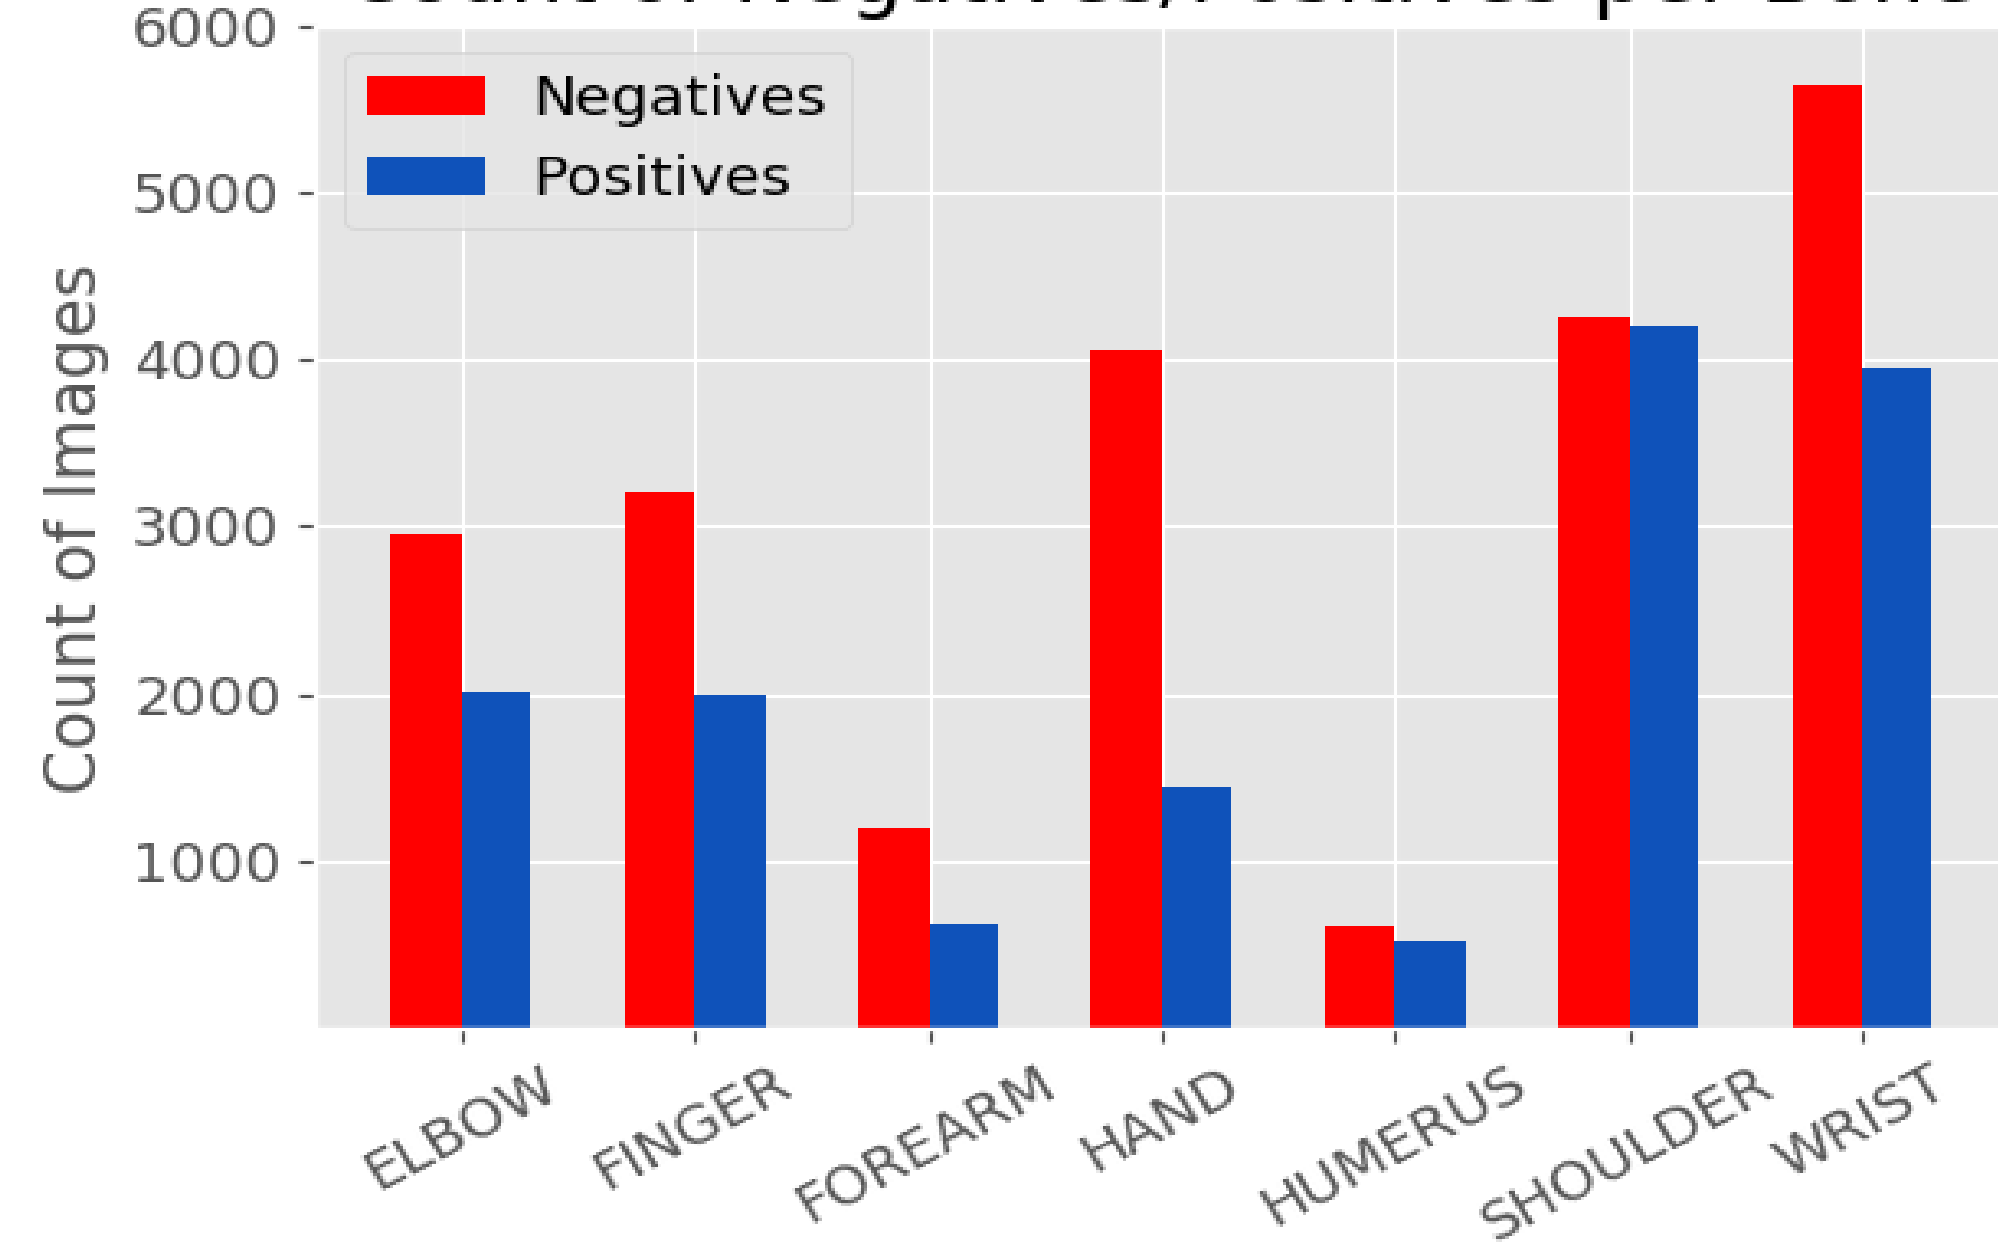
- Test models to make a diagnosis of normal vs abnormal on X-Rays of the arm
 - Simple C – basic CNN
 - CP – convolution-pool
 - CCP - convolution-convolution-pool
- Test best performing model on each set of X-Rays (finger, hand, wrist, forearm, elbow, humerus, shoulder)
- Compare best model performance vs pre-trained model: Xception

Data

- 40k images: 36.5k training images, 3.2k testing images.




Count of Negatives/Positives per Bone



- Difficult data:

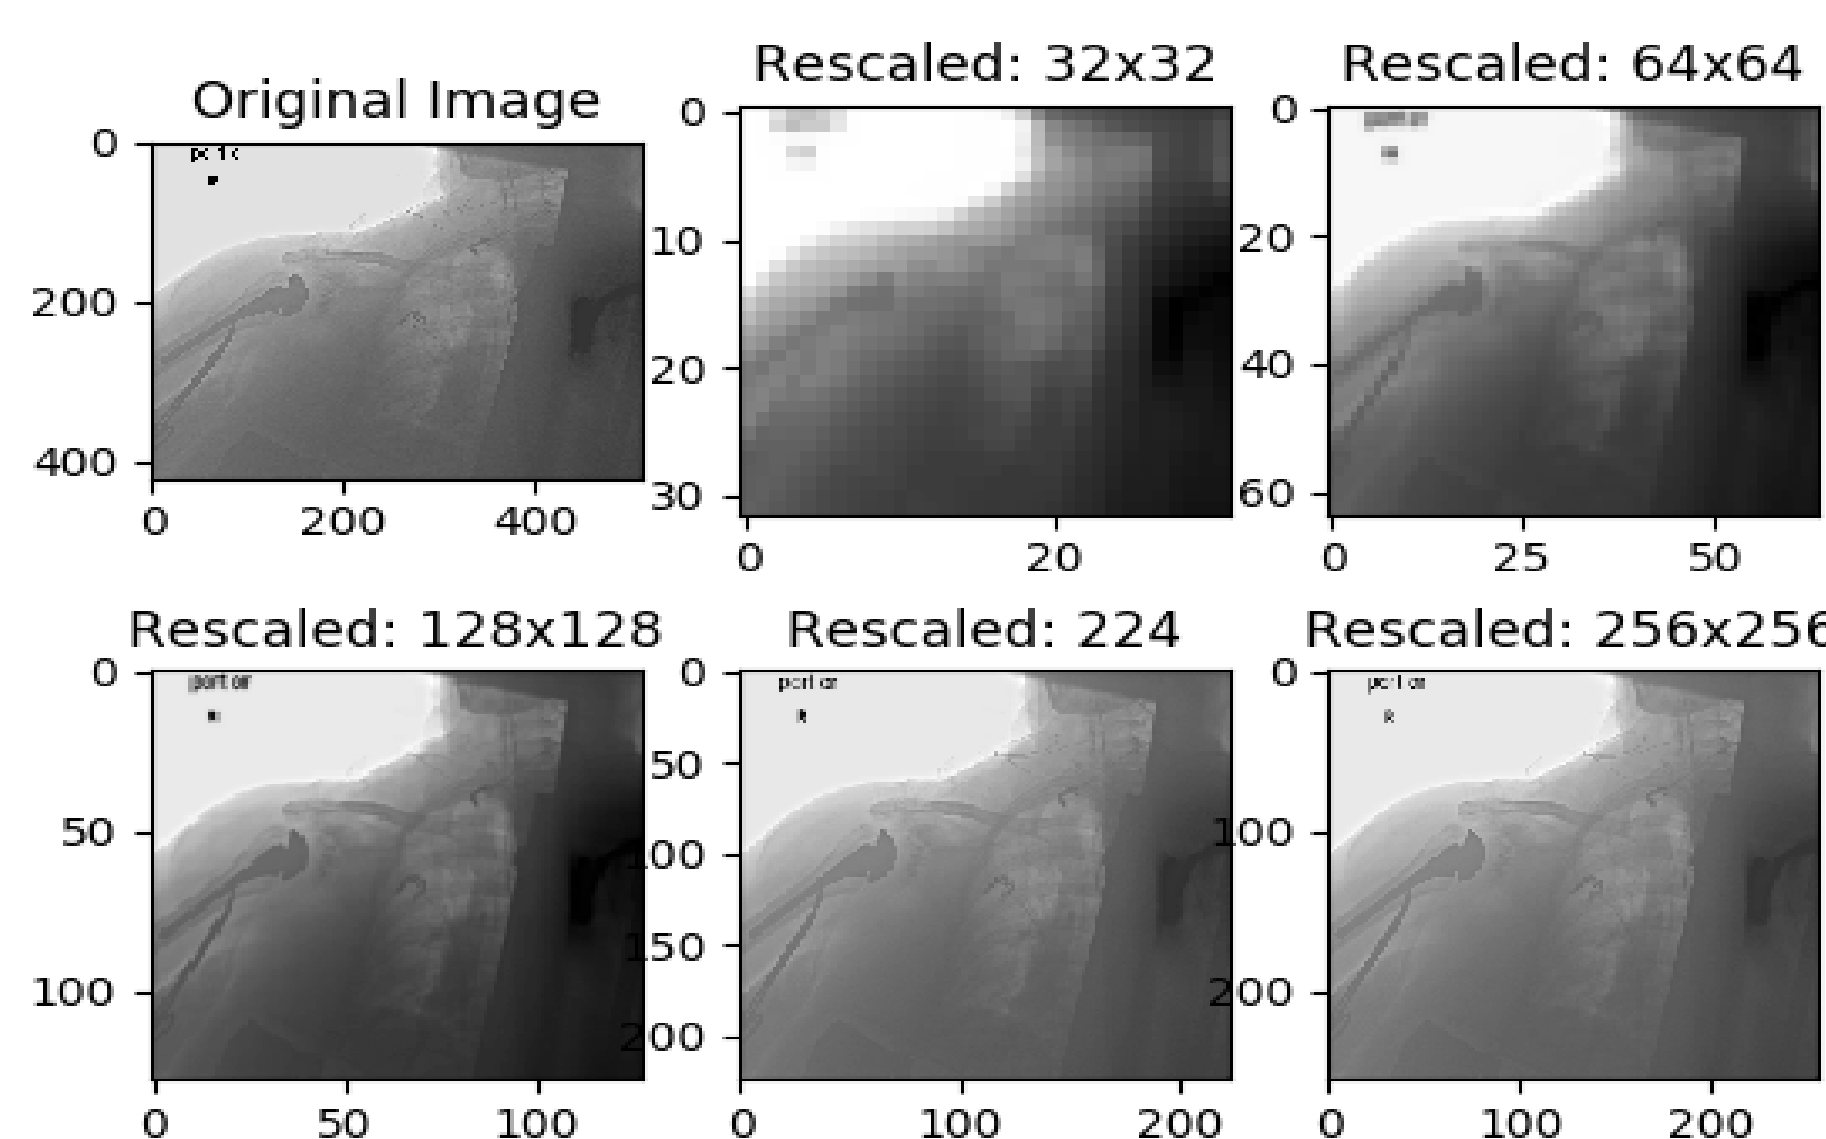


- Misclassified data:



- Total dataset** after cleaning and balancing: 44k train, 3.2k test

Image Size/Augmentation



Results

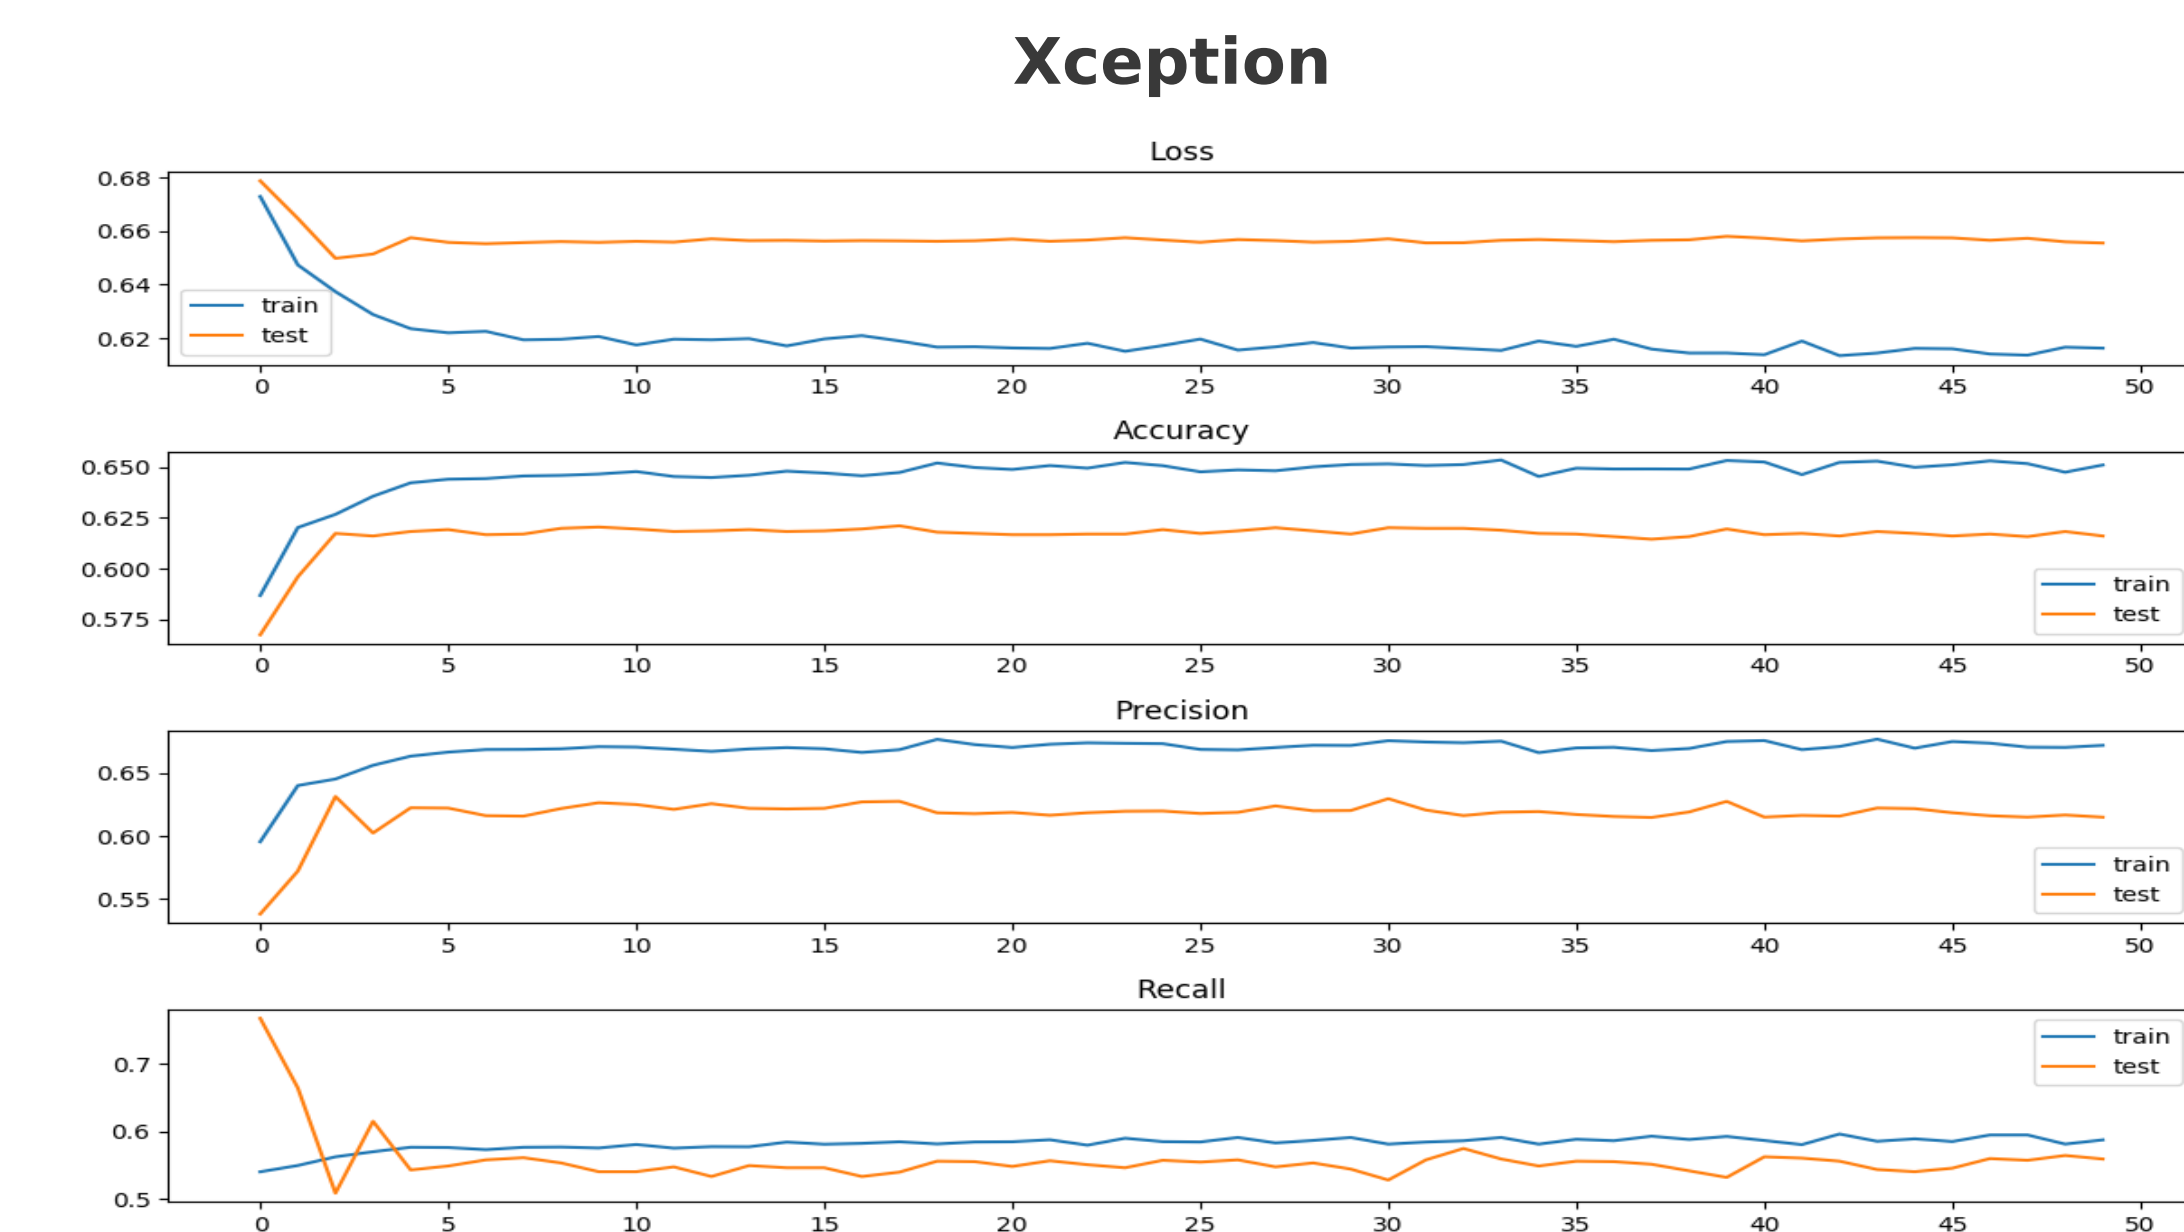



Image Size	64x64		96x96		128x128	
Model	My Model	Xception	My Model	Xception	My Model	Xception
Accuracy	62%	-	63%	62%	61%	63%
Precision	63%	-	65%	62%	61%	67%
Recall	49%	-	58%	56%	56%	45%
F1	55%	-	61%	59%	58%	54%
AUROC	61%	-	63%	61%	61%	62%

F1 Score			Accuracy		
Bone	My Model	Xception	Bone	My Model	Xception
Elbow	73%*	63%	Elbow	75%*	55%
Finger	69%*	64%	Finger	68%*	58%
Forearm	67%*	65%	Forearm	64%*	57%
Hand	21%	56%*	Hand	59%*	54%
Humerus	63%	72%*	Humerus	63%	69%
Shoulder	70%*	12%	Shoulder	73%*	52%
Wrist	66%*	55%	Wrist	70%*	66%

* - indicates better performing model

Model Arch.

Block 1 x 2	
2D Convolutional 3x3 (64 filters)	Activation: Relu/LeakyRelu
2D Max Pooling 2x2	Dropout: 20%

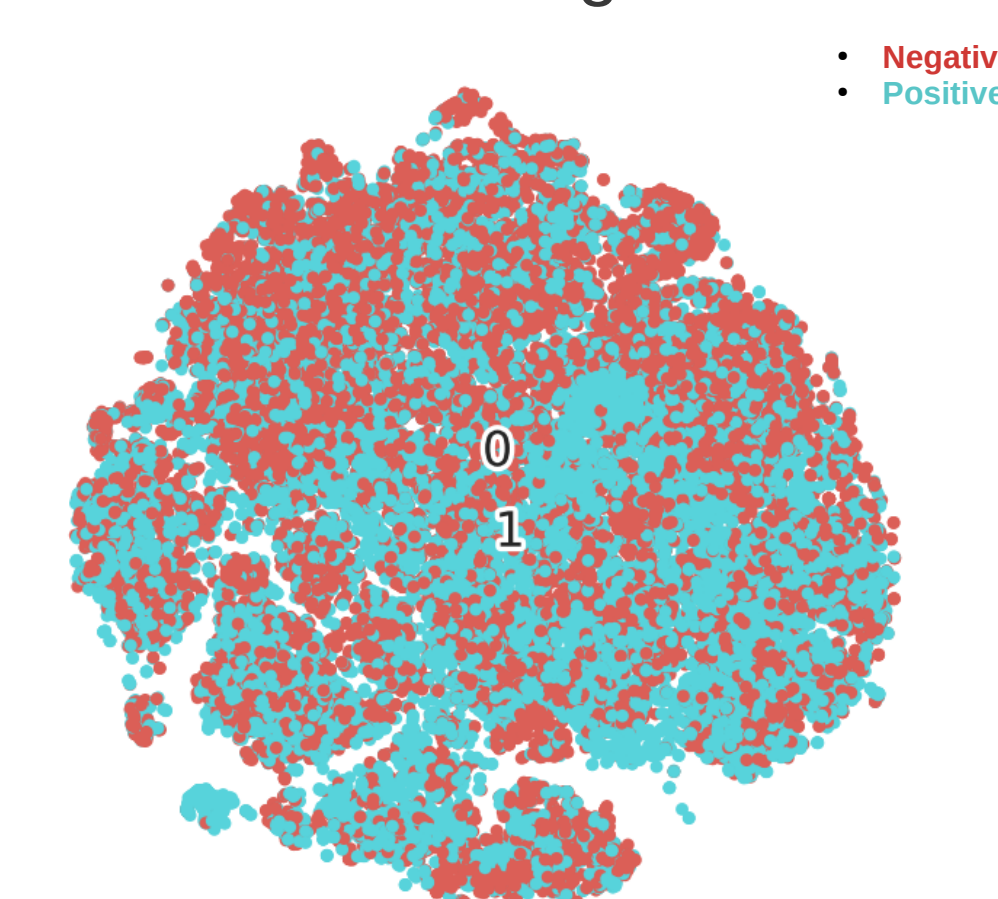
Block 2 x 2	
2D Convolutional 3x3 (128 filters)	Activation: Relu/LeakyRelu
2D Max Pooling 2x2	Dropout: 20%

Block 3 x 2	
2D Convolutional 3x3 (256 filters)	Activation: Relu/LeakyRelu
2D Max Pooling 2x2	Dropout: 20%


Block 4 x 2	
Flatten for Dense Layers	
Dense: 64 Neurons	Activation: LeakyReLU
Dense: 32 Neurons	Activation: LeakyReLU
Dense: 1 Neuron	Activation: Sigmoid
Loss: Binary Crossentropy	
Optimizer: Adam	

Unsupervised


- PCA-TSNE: All Neg. vs Pos.



- PCA-TSNE: All Classes



Misclassified Results



Future Steps

- Separate images
- Visualize layers of CNN
- Class activation map
- Prediction per study
- Adult/youth study
- Transfer Learning
- Unsupervised Learning
- VGG19

References

- Rajpurkar P, Irvin J, Bagu AI, Ding D, Duan T, Mehta H, Yang B, Zhu K, Laird D, Ball RL, et al. 2017. MURA: Large Dataset for Abnormality Detection in Musculoskeletal Radiographs. 1st Conference on Medical Imaging with Deep Learning.

Tech Stack

